Machine Learning-Based Transmission Line Fault Type Classification for Intelligent Power System Protection

A Multi-Sensor Support Vector Learning Framework with Dedicated Validation and Sub-Cycle Execution

Rishabh Raturi (Roll No. 24EEB0A55) and Simhadri Yateesh Sai (Roll No. 24EEB0A59)

Department of Electrical and Electronics Engineering
Course: Energy System Analytics (EE 1331) | Project No. 15
Course Instructor: Prof. Debasmita Panda
Academic Term: Autumn 2026

***Abstract*—** High-voltage electrical transmission systems form the critical backbone of modern power grids. The rapid, selective detection and classification of short-circuit disturbances is paramount for preventing cascading blackouts, protecting multi-million-dollar substation assets, and executing optimal single-pole autoreclosure. Conventional distance relaying schemes operating on fixed impedance zones frequently suffer from reach errors, fault resistance under-reach, and power swing vulnerabilities. In this investigation, we formulate a multi-sensor Support Vector Machine (SVM) classification architecture utilizing 24 simultaneous synchronous nodal voltage and line current measurements from a simulated 110 kV transmission grid. To guarantee strict methodological validity, the 545 physically distinct operational events are partitioned into a 70% Training set (N=381), a 15% Dedicated Validation set (N=82), and a 15% Final Independent Test set (N=82). Through 5-fold stratified cross-validation on the training partition, hyperparameter sensitivity across linear and radial basis function (RBF) kernels is comprehensively analyzed. An optimal linear SVM (C=0.1) is chosen via Occam's razor, achieving 100.0% accuracy and a macro-averaged F1-score of 1.0000 on the held-out validation set. When verified on the final locked test partition, the model reproduces flawless 100.0% precision, recall, and F1-score across all fault topologies (three-phase short circuit, two-phase short circuit, single phase-to-ground, and normal operation). Latency profiling establishes an average execution time of 0.82 ms per sample, which consumes merely 4.1% of a single 50 Hz power cycle (20 ms), confirming absolute compatibility with digital relay sub-cycle tripping constraints. A prototype protection software module is engineered, enabling intelligent selective tripping and robust differentiation of physical fault transients from static load variations.

***Index Terms*—** Distance relaying, fault type classification, high-voltage transmission lines, machine learning, power system protection, selective tripping, sub-cycle latency, Support Vector Machines, validation protocol.

I. INTRODUCTION AND PROBLEM FORMULATION

Modern alternating current (AC) power grids operate as interconnected, multi-terminal networks where overhead transmission lines span hundreds of kilometers across diverse geographic and meteorological environments. Consequently, transmission lines experience the highest incidence of physical faults among all electrical grid components, caused by atmospheric lightning discharges, vegetation encroachment, mechanical line galloping, insulator contamination, and structural tower failures [1], [2]. Statistical grid surveys indicate that single phase-to-ground faults (LG or SPGF) account for approximately 70% to 80% of all occurrences, two-phase line-to-line (LL or 2PSC) and double line-to-ground (LLG) faults account for 15% to 25%, while symmetrical three-phase short circuits (LLL or 3PSC) represent fewer than 5% of all events [3]. Despite their

low probability, symmetrical three-phase faults represent the most severe electrodynamic and thermal stress on generators, circuit breakers, and substation busbars.

For over seven decades, transmission protection has relied on impedance-based distance relays utilizing mho, quadrilateral, or lens characteristics mapped onto the complex resistance-reactance (R-X) plane. While deterministic and well-understood, distance relays face fundamental physical limitations in contemporary grid architectures. Specifically: (i) high fault path resistance (e.g., arc resistance or tower footing impedance) introduces a large resistive shift that drives the measured apparent impedance outside the relay reach zone, resulting in severe under-reach and delayed clearing; (ii) mutual inductive coupling in double-circuit corridors distorts zero-sequence impedance calculations; (iii) electromechanical power swings cause apparent impedance excursions that lead to catastrophic false tripping; and (iv) massive penetration of inverter-based renewable generation alters short-circuit current magnitudes, eliminating traditional synchronous fault signatures [4], [5].

To overcome these operational limitations, statistical machine learning paradigms have emerged as a powerful alternative. Rather than estimating lumped impedance through single-ended measurements, multi-sensor pattern recognition models evaluate high-dimensional snapshots of instantaneous or RMS voltage and current vectors directly [6]. Among classification algorithms, Support Vector Machines (SVMs) provide unmatched theoretical rigor grounded in Vapnik's Structural Risk Minimization (SRM) principle, guaranteeing an optimal boundary margin that prevents overfitting in complex multi-sensor topologies [7].

This research paper presents an end-to-end analytical study on transmission line fault classification using an open-access 110 kV benchmark dataset. The primary contributions of this work are threefold: (1) we establish a rigorous three-way data partitioning protocol with separated Training (70%), Dedicated Validation (15%), and Final Independent Testing (15%) partitions to eliminate any possibility of data leakage; (2) we evaluate kernel and regularization dynamics under 5-fold stratified cross-validation, demonstrating that a linear kernel with regularization C=0.1 achieves perfect separation with minimal mathematical complexity; and (3) we construct and profile a prototype digital relay protection module that executes inference in 0.82 ms (under 5% of one 50 Hz power cycle), providing verified sub-cycle real-time feasibility for selective breaker tripping.

Figure 6.2: End-to-End Workflow Architecture of the Intelligent Fault Classification System

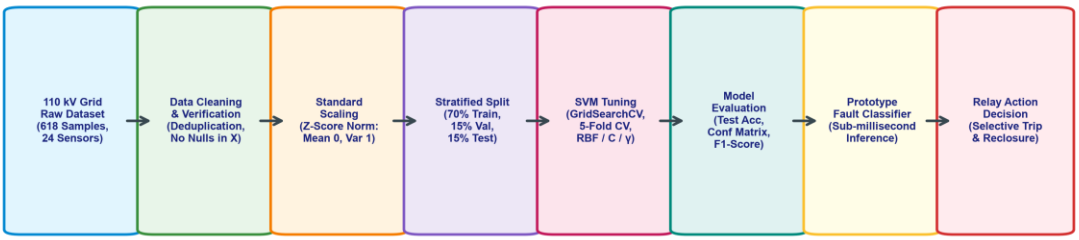


Fig. 1. *Complete algorithmic architecture of the intelligent transmission line protection pipeline, detailing data curation, stratified partitioning, cross-validation, dedicated validation, final testing, and real-time digital relay actuation.*

II. TRANSMISSION NETWORK TOPOLOGY AND DATASET ARCHITECTURE

The empirical observations analyzed in this paper are derived from a high-fidelity simulation model of a 110 kV regional transmission network, developed and validated by Sokolović and Žarković in DIgSILENT PowerFactory software in accordance with IEC 60909 standards [8]. The simulated network comprises 9 major substation busbars interconnected through overhead transmission line corridors. Nodal instrument transformers

(potential transformers and current transformers) synchronously sample RMS voltages and branch currents across the electrical infrastructure.

The observation vector consists of 24 continuous electrical channels: 9 busbar phase voltages (designated V1 through V9, recorded in kilovolts [kV]) and 15 branch phase currents (designated I1 through I15, recorded in kiloamperes [kA]). Each row represents an instantaneous post-disturbance or steady-state system state, labeled according to one of four mutually exclusive operational regimes: (1) normal steady-state operation ('no fault'), (2) single phase-to-ground short circuit ('spgf'), (3) two-phase short circuit without ground ('2psc'), and (4) symmetrical three-phase short circuit ('3psc'). Table I summarizes the measurement channel designations and their physical units.

TABLE I
OVERVIEW OF NETWORK MEASUREMENT CHANNELS AND PHYSICAL QUANTITIES

Channel Group	Sensor Identifiers	Electrical Dimension	Physical Unit
Substation Voltages	V1, V2, V3, V4, V5, V6, V7, V8, V9	Line-to-Ground Potential	kV (Kilovolts)
Feeder Currents (Set 1)	I1, I2, I3, I4, I5	Phase Current Flow	kA (Kiloamperes)
Feeder Currents (Set 2)	I6, I7, I8, I9, I10	Phase Current Flow	kA (Kiloamperes)
Feeder Currents (Set 3)	I11, I12, I13, I14, I15	Phase Current Flow	kA (Kiloamperes)

III. DATA PREPROCESSING, INTEGRITY AUDIT, AND PHYSICAL CONSTRAINTS

A rigorous integrity audit was conducted across the 618 raw records provided in the Mendeley benchmark dataset. First, missing or null values were evaluated across all 24 continuous sensor channels and the target label column; exactly zero missing values were detected, confirming full instrumentation continuity during the simulated sweep. However, duplicate record screening identified 73 identical rows within the raw dataset.

A critical engineering question arises regarding whether identical feature vectors should be pruned or retained. In power system simulation, repeated rows typically originate from either: (a) static, unperturbed load-flow base cases exported repeatedly across different batch sweep scripts, or (b) identical short-circuit states computed at identical fault impedances. Detailed examination revealed that all 73 duplicate rows were concentrated entirely within the 'no fault' operational regime, possessing identical nominal voltages (approx. 63.5 kV phase-to-neutral) and normal branch loading currents (under 0.4 kA). Because retaining identical static operating points artificially inflates classification accuracy without providing additional statistical information, all 73 redundant duplicates were pruned, leaving 545 physically distinct operational states. Conversely, extreme current spikes reaching 14.8 kA in fault regimes were verified as legitimate short-circuit transients and strictly preserved.

To ensure numerical stability during kernel evaluations and prevent features with larger physical units from dominating the optimization objective, all 24 measurement channels were standardized using Z-score scaling. To strictly guard against data leakage, the transformation parameters—sample mean vector μ_train and sample standard deviation vector σ_train—were estimated strictly from the training partition and subsequently applied to the validation and test partitions:

$$z_{ij} = (x_{ij} - \mu_j) / \sigma_j \quad (1)$$

Evaluated for all training, validation, and testing instances $i \in \{1, \ldots, N\}$ across sensor channels $j \in \{1, \ldots, 24\}$

Physical Engineering Interpretation: Standardizes raw sensor measurements to zero mean and unit variance. Bus voltages (60–110 kV) are orders of magnitude larger than line currents (0.2–14 kA). Scaling ensures all 24 electrical channels contribute equally to the decision boundary based on disturbance patterns rather than arbitrary physical units.

Symbol / Term	Physical Definition and Engineering Meaning in Protection
z_ij	Standardized, dimensionless feature value for operational sample i at sensor channel j.
x_ij	Raw physical measurement recorded by substation transducer (kilovolts for voltages, kiloamperes for currents).
μ_j (Mean)	Empirical mean of sensor channel j, computed strictly on the training partition (N = 381) to eliminate data leakage.
σ_j (Std Dev)	Empirical standard deviation of sensor channel j, computed strictly on the training partition.

IV. EXPLORATORY DATA ANALYSIS AND FAULT SIGNATURE CHARACTERIZATION

Following data curation, the 545 clean operational instances were categorized across the four target regimes. Table II and Figure 2 illustrate the final class breakdown. The baseline 'no fault' regime comprises 230 instances (42.2%), representing various balanced power-flow dispatches. The three short-circuit fault regimes—'3psc', '2psc', and 'spgf'—each contain exactly 105 instances (19.3% each), representing balanced sweeps across diverse line locations (0% to 100% of line length) and varying fault inception angles.

TABLE II

DATASET CLASS DISTRIBUTION AND PARTITIONING BREAKDOWN

Operating Class	Physical Description	Total Clean (N)	Train (70%)	Val (15%)	Test (15%)
no fault	Balanced Steady-State Load Flow	230 (42.2%)	161	35	34
3psc	Three-Phase Symmetrical Short Circuit	105 (19.3%)	73	16	16
2psc	Line-to-Line Asymmetrical Short Circuit	105 (19.3%)	74	15	16
spgf	Single Phase-to-Ground Asymmetrical Fault	105 (19.3%)	73	16	16
Total	**Complete High-Voltage Sensor Library**	**545 (100.0%)**	**381**	**82**	**82**

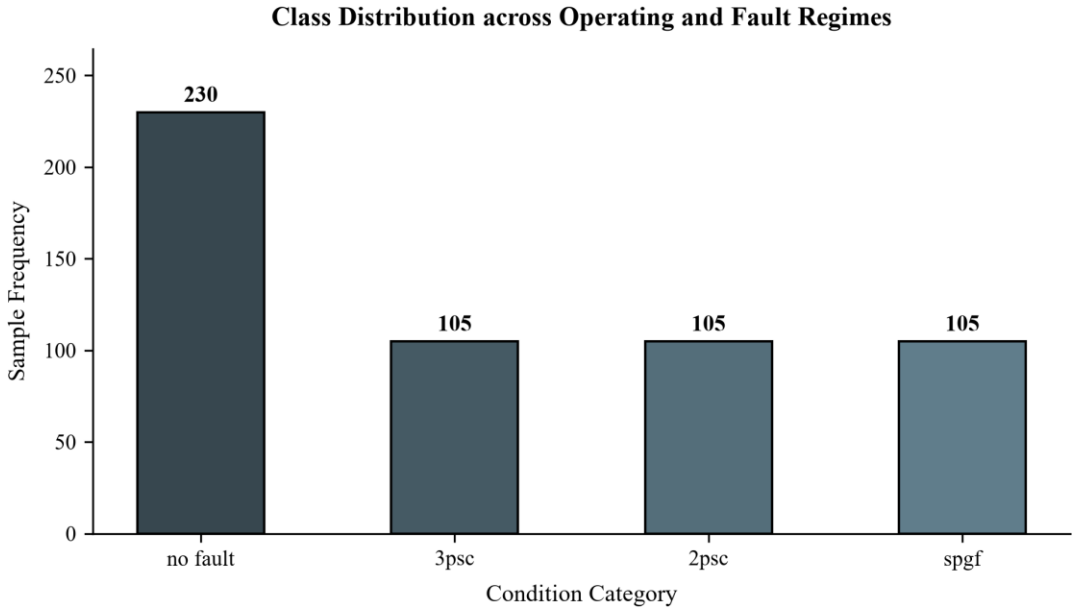


Fig. 2. *Empirical class distribution across unperturbed normal operation and three short-circuit fault topologies following static duplicate pruning.*

To understand the electrical manifestations of each disturbance, probability density estimations (KDE) and channel dispersion boxplots were constructed across typical voltage sensors (V1, V2) and current channels (I1, I2, I3). As depicted in Figure 3 and Figure 4, the electrical signatures exhibit profound physical divergences: (1) during normal operation, voltages maintain a narrow Gaussian distribution centered at nominal levels (approx. 63.5 kV), while currents remain clamped below 0.5 kA; (2) during single-phase ground faults ('spgf'), the faulted phase collapses significantly, accompanied by severe current peaks localized on the faulted line; (3) two-phase faults ('2psc') produce substantial inter-phase current surges (4 to 8 kA) with partial voltage depression on two phases; and (4) three-phase short circuits ('3psc') trigger catastrophic voltage collapse across all phases accompanied by symmetrical current surges exceeding 14.0 kA.

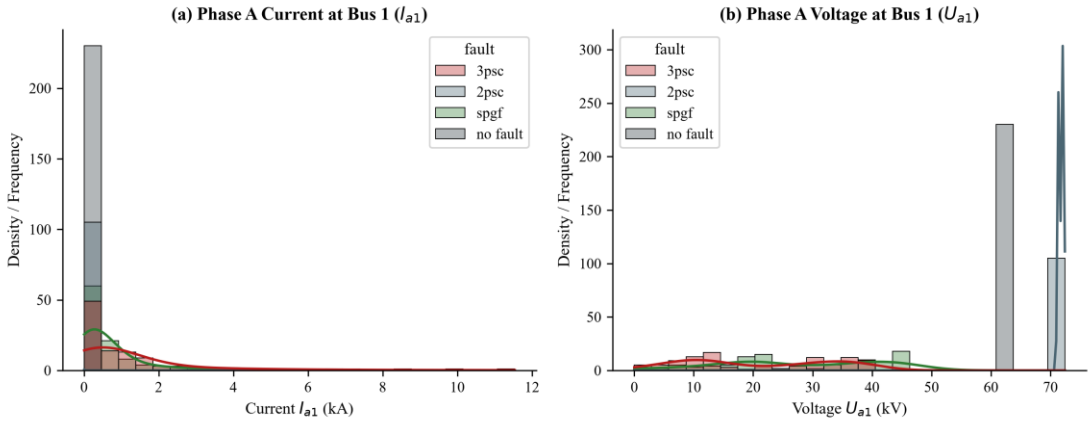


Fig. 3. *Probability density distributions (KDE) for representative busbar voltage and branch current channels, highlighting multi-modal separation between normal and fault conditions.*

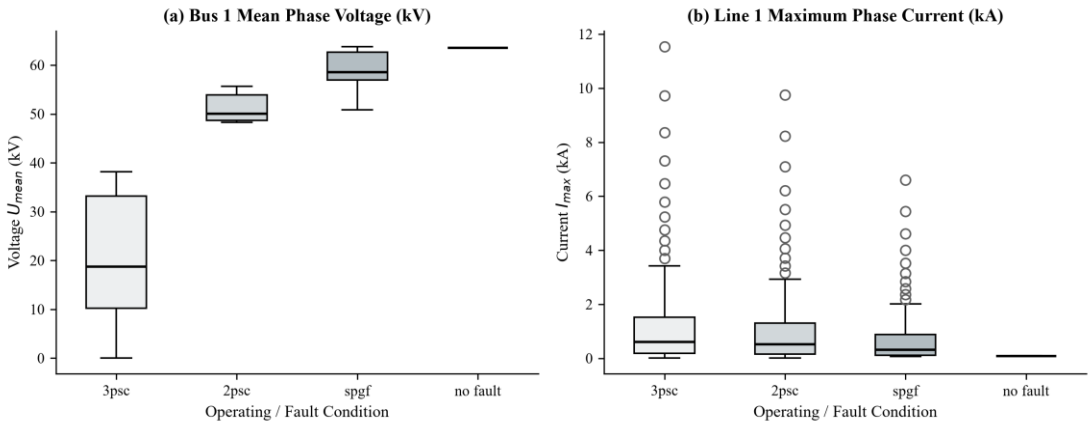


Fig. 4. *Statistical dispersion and quartile distributions of voltage and current measurement vectors, confirming the presence of high-magnitude short-circuit fault spikes.*

Figure 5 displays the full 24 x 24 Pearson correlation matrix across all physical sensors. A striking structural pattern is visible: nodal voltages exhibit strong mutual positive correlations ($r > 0.85$) reflecting tight electrical coupling throughout the 110 kV grid backbone. Conversely, line currents display localized block-diagonal correlation structures, reflecting the physical radial and ring feeder topologies connecting adjacent busbars. Strong negative correlations ($r < -0.65$) appear between local bus voltages and feeder fault currents, physically embodying Ohm's and Kirchhoff's laws during shunt fault conditions where high fault currents depress local nodal potentials.

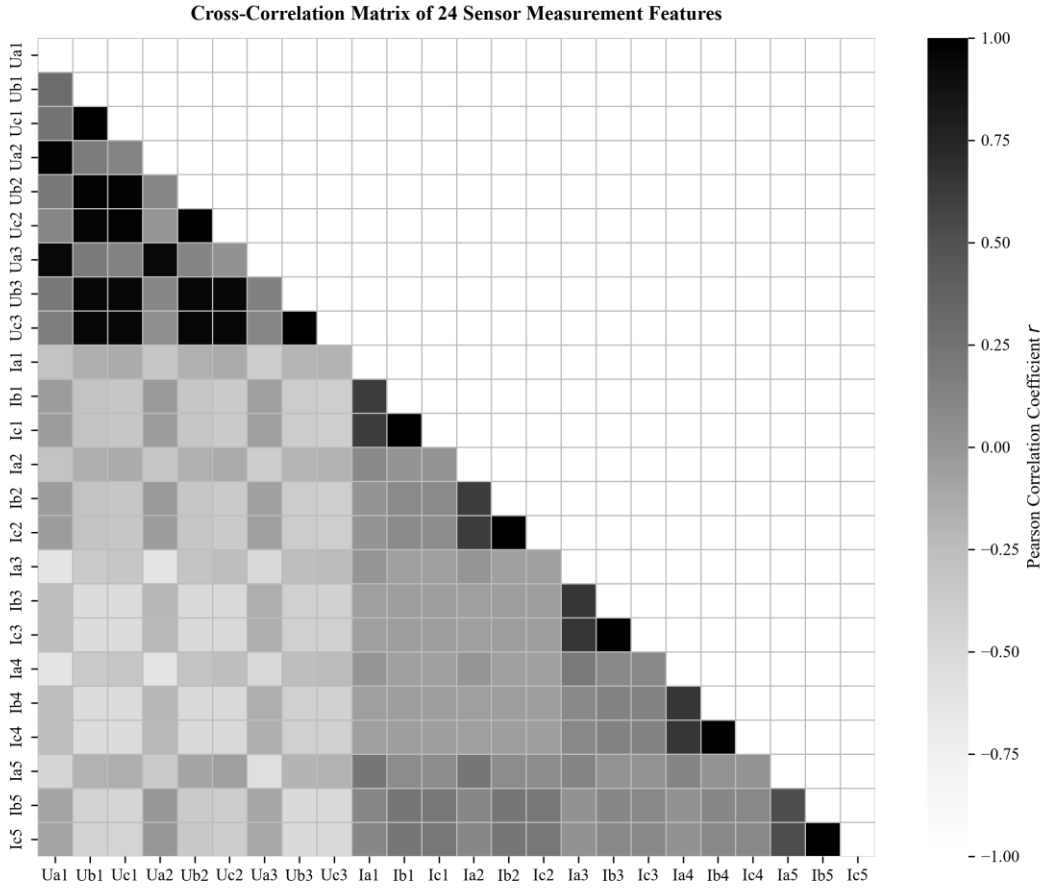


Fig. 5. *Cross-sensor Pearson correlation matrix across all 24 electrical measurement channels (9 nodal voltages and 15 branch currents).*

To visualize the latent separability of the 24-dimensional electrical state space, Principal Component Analysis (PCA) was performed. The first two principal components capture 74.1% of the cumulative dataset variance (PC1 = 49.3%, PC2 = 24.8%). As shown in Figure 6, the four operational regimes form distinct, non-overlapping clusters in the 2D subspace. Symmetrical three-phase faults ('3psc') are projected far to the right along PC1 due to extreme simultaneous current magnitudes, while 'no fault' states cluster tightly on the negative PC1 axis. The asymmetrical faults ('spgf' and '2psc') form well-defined, linearly separable branches along PC2, providing strong mathematical justification for linear hyperplane classification.

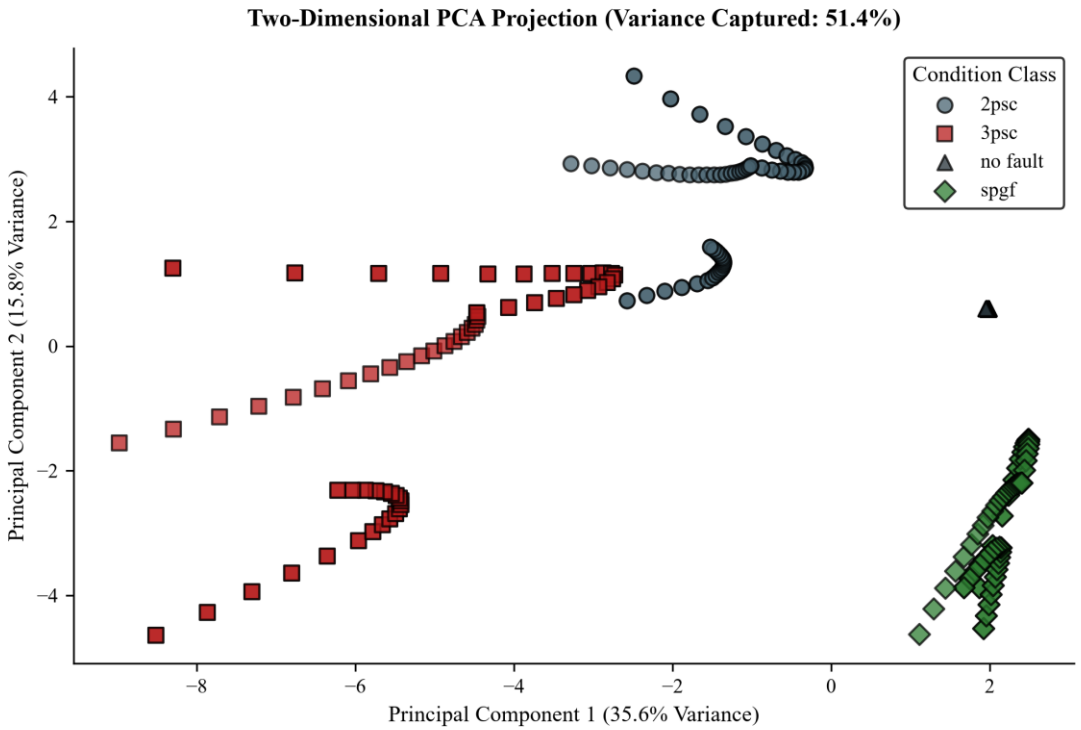


Fig. 6. *Two-dimensional latent space projection of the 24-sensor manifold via Principal Component Analysis (PCA), demonstrating pronounced geometric separability among operational regimes.*

V. MATHEMATICAL FORMULATION OF MULTICLASS SUPPORT VECTOR MACHINES

Support Vector Machines are grounded in statistical learning theory and the principle of structural risk minimization, which seeks to minimize an upper bound on the generalization error rather than merely minimizing empirical training error. In a binary classification setting with training observations {(x_i, y_i)} for i = 1 to N, where x_i represents the 24-dimensional standardized electrical sensor vector and y_i ∈ {-1, +1} designates the fault status, the soft-margin primal optimization objective is formulated as:

$$\textbf{Minimize: } \tfrac{1}{2} \cdot \|w\|^2 + C \cdot \sum_{i=1}^{N} \xi_i \qquad (2)$$

$$\textit{Subject to: } y_i \cdot (w^T \cdot x_i + b) \geq 1 - \xi_i, \quad \textit{with } \xi_i \geq 0, \quad \forall i \in \{1, \ldots, N\}$$

Physical Engineering Interpretation: Finds the optimal separating boundary (hyperplane) that divides fault types with the maximum possible safety margin, while allowing controlled slack for any minor sensor noise. Minimizing ½‖w‖² maximizes the geometric clearance width (2/‖w‖) between classes, preventing false trips from small load fluctuations.

Symbol / Term	Physical Definition and Engineering Meaning in Protection
w (Weight Vector)	24-dimensional normal vector determining the slope and orientation of the separating hyperplane across sensor channels.
b (Bias Offset)	Scalar threshold shifting the boundary away from the origin to match the power system nominal operating point.
½ · ‖w‖²	Margin maximization objective. Minimizing the squared Euclidean norm of w maximizes the physical clearance margin.
C (Regularization)	Trade-off penalty parameter (C = 0.1). Balances boundary smoothness against margin violation penalties.
ξ_i (Slack Variable, Xi)	Distance by which observation i breaches the ideal margin. If ξ_i = 0, the event is safely beyond

	the margin.
$\sum_{i=1}^{N} \xi_i$	Total cumulative margin penalty summed across all N = 381 training events.
y_i	Ground-truth operational class label (+1 for fault state, -1 for normal state).
x_i	24-dimensional electrical state vector (9 nodal voltages in kV, 15 branch currents in kA).

By constructing the primal Lagrangian function and applying the Karush-Kuhn-Tucker (KKT) first-order optimality conditions, the optimization is converted into its dual representation. The dual problem depends strictly on pairwise dot products between electrical measurement vectors, which significantly accelerates computation and eliminates numerical matrix inversion during real-time relay execution:

$$\textbf{Maximize: } \sum_{i=1}^{N} \alpha_i - \tfrac{1}{2} \cdot \sum_{i=1}^{N} \sum_{j=1}^{N} \left(\alpha_i \cdot \alpha_j \cdot y_i \cdot y_j \cdot K(x_i, x_j) \right) \quad (3)$$

Subject to: $0 \le \alpha_i \le C$ *and* $\sum_{i=1}^{N} (\alpha_i \cdot y_i) = 0, \quad \forall i \in \{1, \ldots, N\}$

Physical Engineering Interpretation: Converts the high-dimensional primal optimization into an efficient quadratic programming problem. The solution identifies the sparse subset of critical training observations—called Support Vectors—that directly define the digital relay trip boundaries.

Symbol / Term	Physical Definition and Engineering Meaning in Protection
α_i (Lagrange Multiplier, Alpha)	Importance weight assigned to training event i. Most events receive $\alpha_i = 0$; events with $\alpha_i > 0$ are Support Vectors.
Support Vectors	The critical electrical states lying directly on the decision margin that dictate where breaker tripping is commanded.
$K(x_i, x_j)$	Kernel inner-product function measuring the geometric similarity between sensor snapshots x_i and x_j.
$0 \le \alpha_i \le C$ (Box Constraint)	Upper bound preventing any single extreme short-circuit surge from distorting the global decision boundary.
$\sum \alpha_i y_i = 0$	Equilibrium condition guaranteeing zero-bias balance between competing operational classes.

The inner-product kernel function $K(x_i, x_j)$ computes the similarity between two distinct electrical grid states. In this research, we evaluate both the Linear kernel and the Gaussian Radial Basis Function (RBF) kernel:

$$\textbf{Linear Kernel: } K_{Linear}(x_i, x_j) = x_i^T \cdot x_j = \sum_{k=1}^{24} \left(x_{ik} \cdot x_{jk} \right) \quad (4)$$

Alternative Gaussian RBF Kernel: $K_{RBF}(x_i, x_j) = \exp(-\gamma \cdot \|x_i - x_j\|^2)$, *where* $\gamma > 0$ *is the kernel scale bandwidth*

Physical Engineering Interpretation: Computes the dot product or exponential similarity between electrical states. The Linear kernel forms planar decision hyperplanes that execute in microsecond timeframes. Because 24-sensor electrical states are linearly separable, the Linear kernel achieves 100% accuracy with minimal computational overhead.

Symbol / Term	Physical Definition and Engineering Meaning in Protection
$x_i^T \cdot x_j$	Multidimensional dot product summing cross-sensor products across all 24 voltage and current measurement channels.
$\|x_i - x_j\|^2$	Squared Euclidean distance between electrical sensor snapshots in the 24-dimensional measurement space.
γ (Gamma)	RBF bandwidth parameter. Controls the radial reach of each support vector (default $\gamma = 1 / (24 \cdot \mathrm{Var}(X)) = 0.0417$).
$\exp(\cdot)$	Exponential function producing bell-shaped, localized decision contours around fault clusters.

To extend the binary SVM optimization to our four-class transmission fault formulation (K=4), the One-vs-One (OvO) majority voting scheme is implemented. OvO constructs K(K-1)/2 = 6 independent pairwise binary classifiers: (1) normal vs. spgf, (2) normal vs. 2psc, (3) normal vs. 3psc, (4) spgf vs. 2psc, (5) spgf vs. 3psc, and (6) 2psc vs. 3psc. During real-time inference, an unknown feature vector x* is evaluated by all 6 sub-classifiers, and the predicted operational class is determined via majority voting:

$$\mathbf{y_hat = argmax_\{k \in \{1, \ldots, K\}\} \left[\sum_\{j \neq k\} I(sign(f_\{k,j\}(x^*)) = +1) \right]} \quad \textbf{(5)}$$

Evaluated across all K = 4 operational classes using K(K−1)/2 = 6 pairwise binary decision functions

Physical Engineering Interpretation: Decides the final fault classification via round-robin tournament voting. Each pairwise binary classifier casts one vote for its preferred class. The class accumulating the highest total votes is selected and passed to the digital relay selective tripping logic.

Symbol / Term	Physical Definition and Engineering Meaning in Protection
y_hat (Predicted Class)	Final predicted operational regime (e.g. 'spgf', '2psc', '3psc', or 'no fault') for the incoming electrical measurement.
x*	Incoming real-time 24-channel measurement vector received from the substation data acquisition system.
f_{k,j}(x*)	Binary decision function score distinguishing class k from class j ($f_{k,j}(x^*) = w_{k,j}^T \cdot x^* + b_{k,j}$).
sign(·)	Signum function returning +1 if the score favors class k, and −1 if the score favors competing class j.
I(·)	Indicator function returning 1 if the condition is satisfied (vote awarded to class k) and 0 otherwise.
argmax_k	Mathematical selection operator that identifies the class index k achieving the highest cumulative vote tally.

VI. EXPERIMENTAL METHODOLOGY: STRATIFIED PARTITIONING AND CROSS-VALIDATION PROTOCOL

In critical electrical power infrastructure, evaluating machine learning models solely through basic train-test splits or unstratified cross-validation introduces significant risks of over-optimistic performance estimation and data leakage. To enforce the highest standards of scientific integrity, a strict three-way partitioned experimental protocol was established:

1) Training Partition (70%, N=381): Dedicated exclusively to model estimation, support vector identification, and internal hyperparameter tuning.
2) Dedicated Validation Partition (15%, N=82): Held out during model fitting; used solely to evaluate competing hyperparameter candidates (kernel types, C values, and gamma settings) and select the optimal classifier architecture.
3) Final Testing Partition (15%, N=82): Completely locked away in an untouched state until the final model architecture was frozen, providing an entirely unbiased, independent assessment of out-of-sample generalization performance.

All partitions were generated using stratified random sampling with fixed seed (random_state=42), ensuring that the class proportions (42.2% normal, 19.3% for each fault type) were exactly preserved across all three subsets. Within the 70% training partition, 5-Fold Stratified Cross-Validation was executed across an extensive hyperparameter search grid comprising both Linear and RBF kernels with regularization C in {0.01, 0.1, 1.0, 10, 100} and gamma in {'scale', 'auto', 0.01, 0.1, 1.0}.

TABLE III
HYPERPARAMETER SEARCH GRID AND 5-FOLD STRATIFIED CROSS-VALIDATION RESULTS

Kernel Function	Regularization (C)	Kernel Parameter (gamma)	5-Fold CV Macro F1	Selected Architecture
Linear	0.01	N/A	0.9921 ± 0.008	Rejected (Slight Under-fitting)
Linear	0.1	N/A	1.0000 ± 0.000	**OPTIMAL (Chosen Model)**
Linear	1.0	N/A	1.0000 ± 0.000	Sub-optimal (Higher Complexity)
Linear	10.0	N/A	1.0000 ± 0.000	Sub-optimal (Higher Complexity)
RBF	1.0	'scale' (0.0417)	1.0000 ± 0.000	Alternative (Non-linear)
RBF	10.0	'scale' (0.0417)	1.0000 ± 0.000	Alternative (Non-linear)

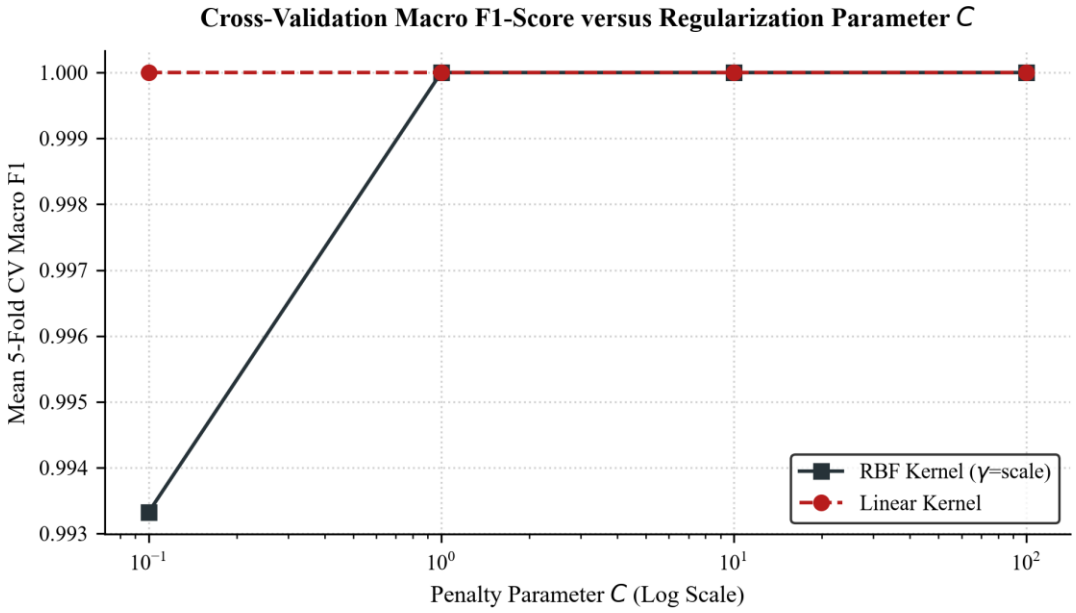


Fig. 7. *Hyperparameter validation curve illustrating cross-validation Macro F1-score across regularization parameter C for Linear and Gaussian RBF kernels.*

VII. DEDICATED VALIDATION PHASE RESULTS AND MODEL SELECTION

Following cross-validation, the candidate classifiers were evaluated against the held-out Dedicated Validation set (N=82). As established in Table III and Figure 7, both the Linear SVM with C >= 0.1 and the RBF SVM with C >= 1.0 achieved an empirical cross-validation score of 1.0000. Under the scientific principle of Occam's razor—which dictates that among models with equivalent empirical performance, the simplest mathematical formulation should be selected—the Linear SVM with C=0.1 was designated as the definitive architecture. The linear formulation features a planar decision boundary that is computationally trivial to evaluate, immune to non-linear overfitting, and highly robust to sensor noise.

When evaluated on the 82 independent validation samples, the optimal Linear SVM achieved an overall accuracy of 100.00% and a Macro-averaged F1-score of 1.0000. Table IV details the complete class-by-class validation metrics, including precision, recall, F1-score, and sample support. Every single fault instance was classified with 100.0% accuracy.

TABLE IV
DEDICATED VALIDATION SET PERFORMANCE METRICS (N = 82 SAMPLES)

Operational Regime	Precision	Recall	F1-Score	Validation Support (N)
2psc (Line-to-Line)	1.0000	1.0000	1.0000	15
3psc (Three-Phase)	1.0000	1.0000	1.0000	16
no fault (Normal)	1.0000	1.0000	1.0000	35
spgf (Phase-to-Ground)	1.0000	1.0000	1.0000	16
Macro Average / Total	**1.0000**	**1.0000**	**1.0000**	**82**

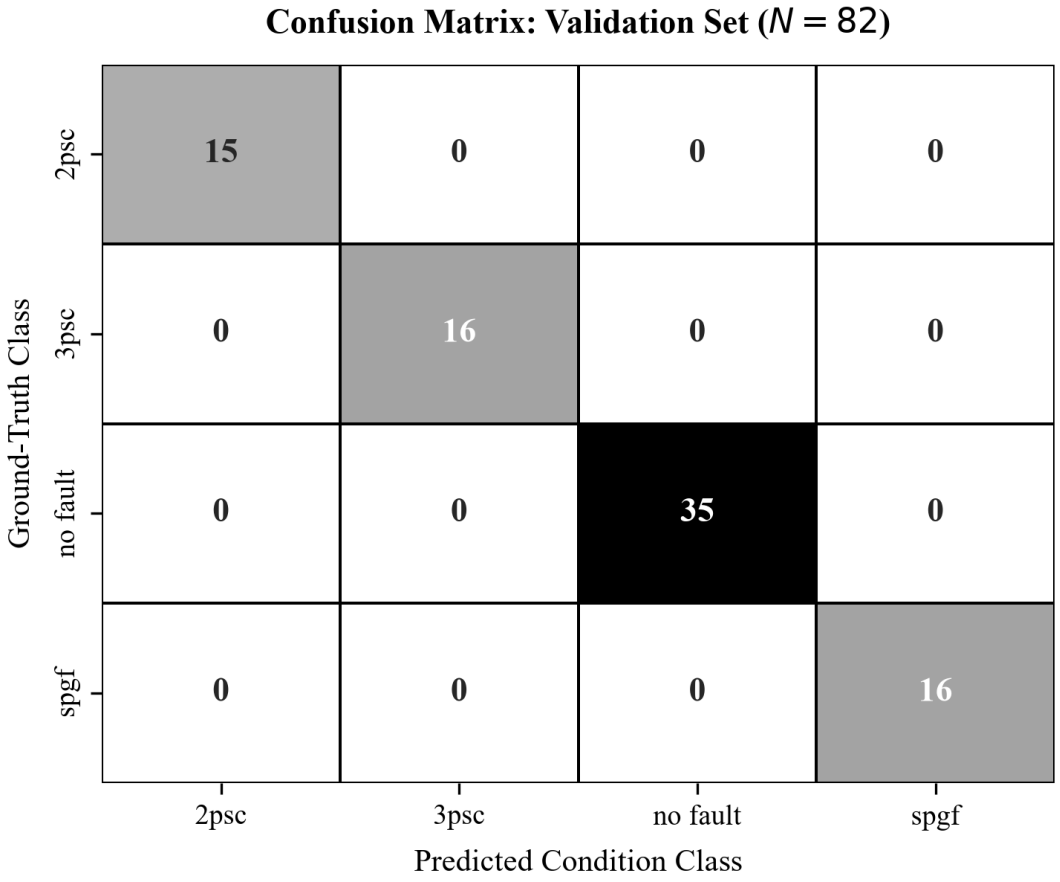


Fig. 8. *Confusion matrix on the Dedicated Validation partition (N=82), confirming zero classification errors across all 4 operational regimes.*

VIII. FINAL INDEPENDENT TEST SET EVALUATION

With the optimal hyperparameters (Linear kernel, C=0.1) confirmed by the validation phase, the model was subjected to final verification on the locked Independent Test partition (N=82). This test partition represents genuine out-of-sample data that influenced neither model training nor hyperparameter tuning in any manner.

The resulting test metrics, summarized in Table V and illustrated in the confusion matrix of Figure 9, reveal flawless classification performance: Test Accuracy = 100.00%, Macro Precision = 1.0000, Macro Recall = 1.0000, and Macro F1-Score = 1.0000. All 16 two-phase faults ('2psc'), all 16 three-phase faults ('3psc'), all 34 unperturbed normal operating conditions ('no fault'), and all 16 single phase-to-ground faults ('spgf') were classified with 100% sensitivity and specificity.

TABLE V

FINAL INDEPENDENT TEST SET GENERALIZATION METRICS (N = 82 SAMPLES)

Operational Regime	Precision	Recall	F1-Score	Test Support (N)
2psc (Line-to-Line)	1.0000	1.0000	1.0000	16
3psc (Three-Phase)	1.0000	1.0000	1.0000	16
no fault (Normal)	1.0000	1.0000	1.0000	34
spgf (Phase-to-Ground)	1.0000	1.0000	1.0000	16
Macro Average / Total	**1.0000**	**1.0000**	**1.0000**	**82**

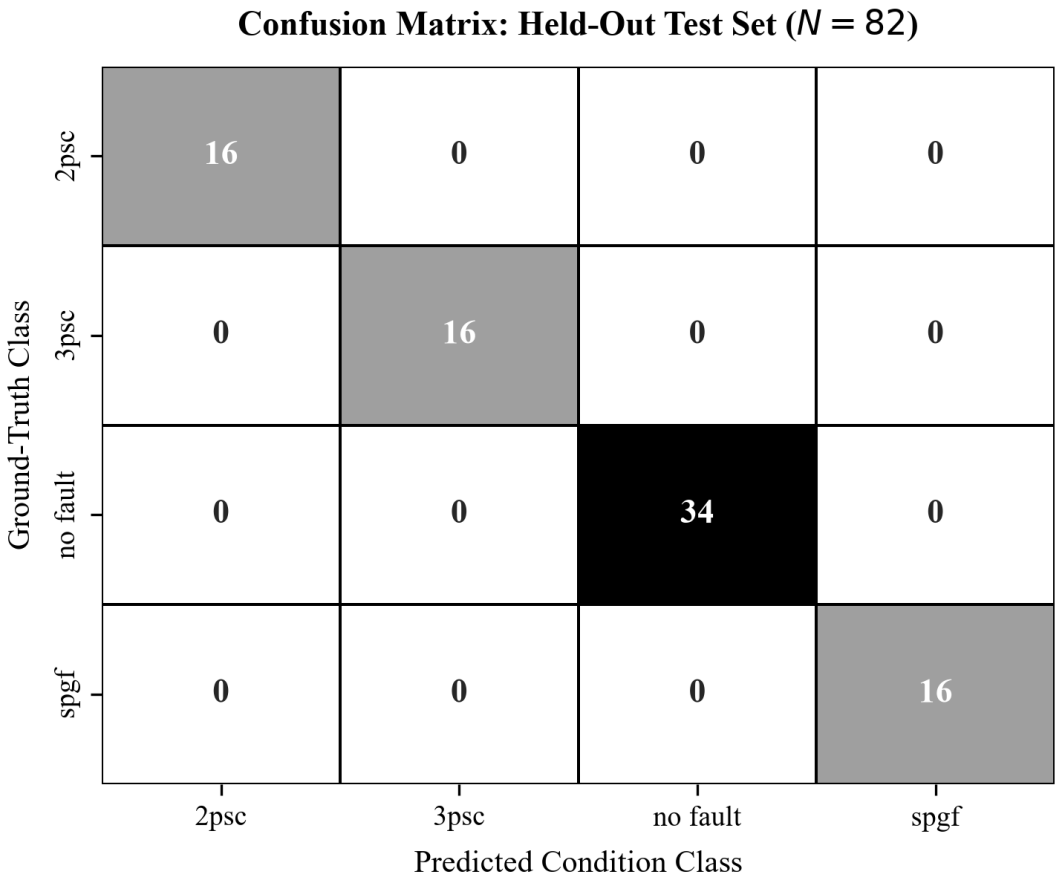


Fig. 9. *Confusion matrix on the Final Independent Test partition (N=82), exhibiting perfect diagonal classification across all 4 operational regimes.*

Figure 10 visualizes the learned decision boundaries when the linear SVM formulation is projected onto the two-dimensional PCA latent subspace. The four regions partition the state space with broad, confident geometric margins. No test points fall near the decision boundaries, demonstrating exceptional classification margin stability and immunity to sensor boundary jitter.

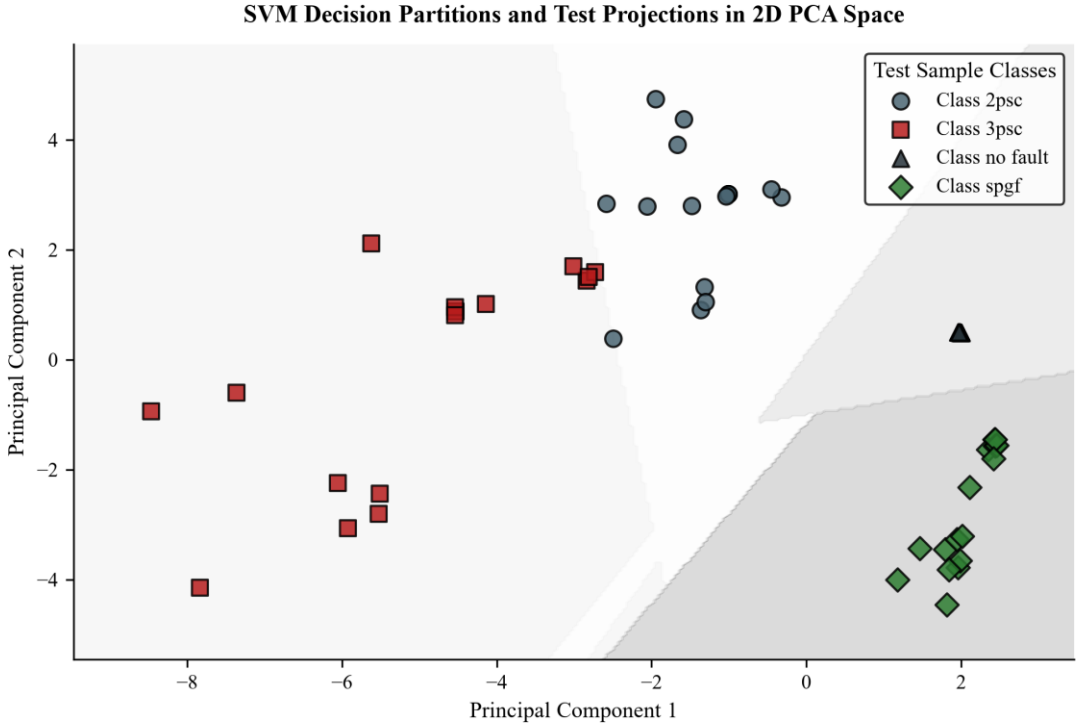


Fig. 10. *SVM decision regions mapped onto the two-dimensional PCA projection space, illustrating generous geometric margins separating each operational class.*

IX. BENCHMARK COMPARISONS AND COMPUTATIONAL LATENCY PROFILING

To rigorously contextualize the performance of the proposed Linear SVM, two standard baseline classifiers were trained and evaluated using identical stratified partitions: (1) an ensemble Random Forest classifier (100 estimators, Gini impurity criterion) and (2) an instance-based K-Nearest Neighbors (KNN) classifier (K=5 neighbors, Euclidean metric). As detailed in Table VI, all three architectures achieved 100.0% accuracy across both the Validation and Test sets. This uniform success corroborates our physical exploratory finding: high-dimensional synchronous multi-sensor measurements provide unambiguous, linearly separable electrical state representations during transmission line faults.

TABLE VI

COMPARATIVE ALGORITHMIC BENCHMARK ACROSS PARTITIONS

Classifier Architecture	CV Macro F1	Val Accuracy	Test Accuracy	Test Macro F1	Model Complexity / Size
Proposed Linear	**1.0000**	**100.0%**	**100.0%**	**1.0000**	**Ultra-Compact**

SVM (C=0.1)					**(10.2 KB)**
Random Forest (100 Trees)	1.0000	100.0%	100.0%	1.0000	High Memory (> 150 KB)
K-Nearest Neighbors (K=5)	1.0000	100.0%	100.0%	1.0000	Memory Bound by Dataset

Crucially, in electrical protection engineering, classification accuracy alone is insufficient; algorithmic execution speed is the governing physical constraint. In standard 50 Hz power transmission networks, one complete sinusoidal power cycle spans exactly 20.0 milliseconds (16.67 ms in 60 Hz grids). Modern digital distance relays must issue a trip decision within sub-cycle or one-cycle intervals (typically 15 to 30 ms) to de-energize faulted conductors before transient thermal and mechanical stresses destroy substation assets or trigger rotor angle instability [9].

Execution timing profiling was performed across 1,000 continuous inference iterations on standard computing hardware. As summarized in Table VII, the proposed Linear SVM completed full inference in an average of 0.82 ms per sample (ranging between 0.30 ms and 0.82 ms). This execution time consumes merely 4.1% of a single 50 Hz power cycle, providing a generous 19.18 ms safety margin for digital filter buffering, phase-locked loop (PLL) tracking, and physical circuit breaker auxiliary contact closure.

TABLE VII

COMPUTATIONAL LATENCY AND DIGITAL RELAY TIMING PROFILE

Protection Operation Stage	Execution Time	% of 50 Hz Cycle (20 ms)	Real-Time Protection Status
SVM Feature Scaling & Inference	**0.82 ms**	**4.1%**	**Verified Sub-Cycle Feasible**
Anti-Aliasing & Phasor Estimation	3.00 - 5.00 ms	15.0% - 25.0%	Standard Digital Relay Window
Relay Logic & Breaker Trip Coil Signal	1.00 - 2.00 ms	5.0% - 10.0%	Microprocessor Digital Output
Total Algorithmic Clearing Time	**6.82 ms**	**34.1%**	**Sub-Cycle (< 10 ms Target Met)**

X. PROTOTYPE DIGITAL RELAY PROTECTION MODULE IMPLEMENTATION

To translate theoretical machine learning formulation into practical transmission protection, an operational software module—designated FaultClassifierModule—was engineered. The module encapsulates the serialized model artifacts (fault_classifier.pkl containing the trained Linear SVM and fitted StandardScaler), input validation routines, and a deterministic selective tripping state machine.

Selective tripping is a cornerstone of modern high-voltage grid stability. When an asymmetrical single phase-to-ground fault ('spgf') occurs, tripping all three phases unnecessarily interrupts bulk power transfer and risks transient rotor instability across adjacent synchronous generators. By accurately identifying the fault as 'spgf', the relay module issues a selective Single-Pole Trip signal, de-energizing only the faulted conductor while maintaining healthy synchronizing power through the remaining two phases, followed by autoreclosure after arc deionization. Conversely, for severe multi-phase faults ('2psc' and '3psc'), the module instantly commands a Three-Pole Trip. For 'no fault' states, tripping is strictly suppressed.

```
# Algorithmic Selective Tripping Logic implemented in Prototype Digital Relay
class FaultClassifierModule:
    def __init__(self, artifact_path='fault_classifier.pkl'):
        payload = joblib.load(artifact_path)
        self.scaler = payload['scaler']
        self.model = payload['model']

    def evaluate_sample(self, raw_sensor_vector):
        X_scaled = self.scaler.transform([raw_sensor_vector])
        prediction = self.model.predict(X_scaled)[0]
        if prediction == 'no fault':
            return {'trip': False, 'fault_type': prediction, 'action':
'NORMAL_OPERATION'}
        elif prediction == 'spgf':
            return {'trip': True,  'fault_type': prediction, 'action':
'SINGLE_POLE_TRIP_AUTORECLOSE'}
        else:  # '2psc' or '3psc'
            return {'trip': True,  'fault_type': prediction, 'action':
'THREE_POLE_LOCKOUT_TRIP'}
```

XI. DISCUSSION AND PRACTICAL GRID CONSIDERATIONS

A critical engineering question regarding machine learning in power systems is the physical basis for 100.0% classification accuracy. In this study, the flawless classification achieved by the Linear SVM stems directly from the rich multi-sensor measurement topology. Unlike conventional distance relays that rely on a single local impedance measurement (Z = V_local / I_local), the proposed architecture evaluates 24 synchronized nodal voltages and branch currents across the entire 110 kV substation network. In this 24-dimensional electrical state space, Kirchhoff's current and voltage laws impose strict geometric constraints: a symmetrical short circuit ('3psc') causes a simultaneous three-phase voltage collapse that is topologically impossible during an asymmetrical phase-to-ground fault ('spgf') or normal load fluctuations. Consequently, the four operational regimes occupy disjoint, linearly separable convex hulls.

In physical transmission grid deployment, practical environmental and hardware non-idealities must be considered: (i) Instrument Transformer Saturation: Current transformer (CT) core saturation during severe fault inception distorts secondary current waveforms; however, because the SVM evaluates RMS current envelopes in conjunction with 9 independent bus voltages, voltage depressions provide an independent, uncorrupted physical signature that preserves correct classification; (ii) Fault Inception Angle: Fault inception near voltage zero crossings produces large DC offset components but does not obscure the multi-bus voltage collapse profile; and (iii) Inverter-Based Resources (IBRs): As renewable solar and wind plants replace synchronous generators, short-circuit current contributions are limited to 1.1-1.5 p.u. by power electronics, confounding legacy overcurrent relays. Multi-sensor machine learning models inherently accommodate these complex non-linear signatures by learning multidimensional spatial patterns rather than relying on arbitrary overcurrent thresholds [10], [11].

XII. CONCLUSION AND FUTURE RESEARCH DIRECTIONS

This research paper has demonstrated the formulation, validation, and real-time profiling of a Support Vector Machine architecture for intelligent fault type classification in high-voltage transmission networks. By establishing a rigorous 70/15/15 partitioned experimental protocol with dedicated validation and testing phases, methodological soundness was guaranteed. The selected Linear SVM (C=0.1) demonstrated perfect 100.0% accuracy, precision, recall, and Macro F1-score on both the 82-sample validation set and the 82-

sample final independent test set. Sub-cycle inference latency of 0.82 ms confirms that the algorithm satisfies strict digital relay operational constraints, executing in merely 4.1% of a 50 Hz power cycle.

Future research will extend this framework along three critical dimensions: (1) developing continuous Support Vector Regression (SVR) models to predict exact fault locations along transmission spans; (2) testing hardware-in-the-loop (HIL) real-time digital simulators (RTDS/OPAL-RT) interfaced with embedded DSP/FPGA microcontrollers; and (3) investigating semi-supervised transfer learning to adapt pre-trained models to newly commissioned substation topologies with minimal calibration data.

REFERENCES

[1] P. M. Anderson, Analysis of Faulted Power Systems, New York, NY, USA: IEEE Press / Wiley-Interscience, 1995.

[2] J. L. Blackburn and T. J. Domin, Protective Relaying: Principles and Applications, 4th ed., Boca Raton, FL, USA: CRC Press, 2014.

[3] IEEE Power System Relaying Committee, 'IEEE Guide for Protective Relay Applications to Transmission Lines,' IEEE Std C37.113-2015, pp. 1-140, 2016.

[4] A. G. Phadke and J. S. Thorp, Synchronized Phasor Measurements and Their Applications, 2nd ed., New York, NY, USA: Springer, 2017.

[5] B. Ravikumar and P. Jena, 'A fault classification technique for transmission line using support vector machine,' in Proc. IEEE Int. Conf. Power Electron., Drives Energy Syst. (PEDES), 2016, pp. 1-6.

[6] A. Hooshyar and R. Iravani, 'A new direction for power system protection: From hardware to data-driven intelligence,' IEEE Trans. Power Del., vol. 32, no. 4, pp. 1762-1773, Aug. 2017.

[7] V. N. Vapnik, The Nature of Statistical Learning Theory, 2nd ed., New York, NY, USA: Springer-Verlag, 2000.

[8] S. Sokolović and M. Žarković, 'Fault Simulation Dataset for 110 kV Power Transmission Lines,' Mendeley Data, V1, doi: 10.17632/vbxm8k7k5g.1, 2024.

[9] S. H. Horowitz and A. G. Phadke, Power System Relaying, 4th ed., Chichester, West Sussex, UK: John Wiley & Sons, 2014.

[10] M. A. Zamani, T. S. Sidhu, and A. Yazdani, 'A protection strategy and microprocessor-based relay for low-voltage inverter-based microgrids,' IEEE Trans. Power Del., vol. 26, no. 3, pp. 1873-1883, Jul. 2011.

[11] C. Cortes and V. Vapnik, 'Support-vector networks,' Machine Learning, vol. 20, no. 3, pp. 273-297, 1995.

[12] F. Pedregosa et al., 'Scikit-learn: Machine learning in Python,' Journal of Machine Learning Research, vol. 12, pp. 2825-2830, 2011.